

Customer Churn Prediction Using Machine Learning

Monika Tyagi
Indian Institute of Science
Email: monikatyagi@iisc.ac.in

Sourajit Bhar
Indian Institute of Science
Email: sourajitbhar@iisc.ac.in

Abstract—Customer churn prediction is a critical task for subscription-based businesses. In this project, we analyze the Telco Customer Churn dataset and apply multiple supervised machine learning models—including Logistic Regression, Decision Tree, Random Forest, Naive Bayes, and Support Vector Machine (SVM)—to detect at-risk customers. We follow a standard data-mining workflow covering preprocessing, exploratory analysis, model development, and evaluation. Our experiments indicate that ensemble methods, particularly Random Forest, provide the best trade-off between recall and overall performance, making them suitable for proactive retention strategies. Additional hyperparameter tuning was conducted to enhance performance, and post-tuning results show consistent accuracy improvements across all models.

Index Terms—Customer Churn, Classification, Data Mining, Machine Learning, Random Forest, Hyperparameter Tuning, Telco

I. INTRODUCTION

Customer churn refers to the phenomenon of customers discontinuing a service. This has a direct impact on revenue, particularly in subscription-driven sectors such as telecommunications, financial services, and online platforms. Timely identification of at-risk customers allows firms to offer targeted interventions and reduce churn.

This work applies core techniques from DA 2270 to develop churn prediction models on a widely used public dataset. Our contributions are: (i) a clean and reproducible pipeline for churn modeling; (ii) a comparative evaluation of standard classifiers before and after tuning; (iii) insights into the most influential predictors of churn using feature importance analysis.

II. RELATED WORK

Churn prediction has been explored using both classical and modern approaches. Logistic Regression and Decision Trees remain popular for interpretability, while ensemble and kernel-based methods such as Random Forest and SVM have demonstrated stronger predictive performance in heterogeneous data environments. Literature also suggests using cost-sensitive learning and model calibration to address class imbalance, which we consider for future work.

III. DATASET

We use the *Telco Customer Churn* dataset (Kaggle), containing 7,043 customer records with demographic details, service usage, contract type, payment method, and churn labels. The

TABLE I: Model Performance (Before Tuning)

Model	Acc.	Prec.	Recall	F1
Logistic Regression	0.800	0.660	0.620	0.640
Decision Tree	0.790	0.610	0.650	0.630
Random Forest	0.830	0.700	0.680	0.690
SVM (RBF)	0.810	0.710	0.600	0.650
Naive Bayes	0.780	0.600	0.580	0.590

target variable is binary: Churn = Yes/No. Data preprocessing included:

- Removal of redundant columns (e.g., customerID)
- Handling of missing values in TotalCharges
- One-hot encoding of categorical features
- Normalization of continuous attributes such as tenure and MonthlyCharges

IV. METHODOLOGY

Our modeling pipeline follows a structured workflow:

- 1) **Data Cleaning and Preprocessing:** Data consistency checks, imputation, and feature encoding.
- 2) **Exploratory Data Analysis (EDA):** Distribution plots, correlation matrices, and churn percentage visualization.
- 3) **Feature Engineering:** Derived binary flags for service bundles and total monthly cost.
- 4) **Model Training:** Logistic Regression, Decision Tree, Random Forest, Naive Bayes, and SVM (RBF kernel).
- 5) **Evaluation:** Metrics include accuracy, precision, recall, F1-score, ROC–AUC, and confusion matrices.
- 6) **Hyperparameter Tuning:** Conducted using GridSearchCV and cross-validation for optimal regularization (C), tree depth, number of estimators, and kernel parameters.

V. EXPERIMENTS AND RESULTS

- A. Initial Model Comparison
- B. Confusion Matrices (Before Tuning)
- C. Model Optimization and Hyperparameter Tuning

After grid search tuning, all models improved in accuracy and recall. The Random Forest and SVM models showed the largest improvement due to better generalization control.

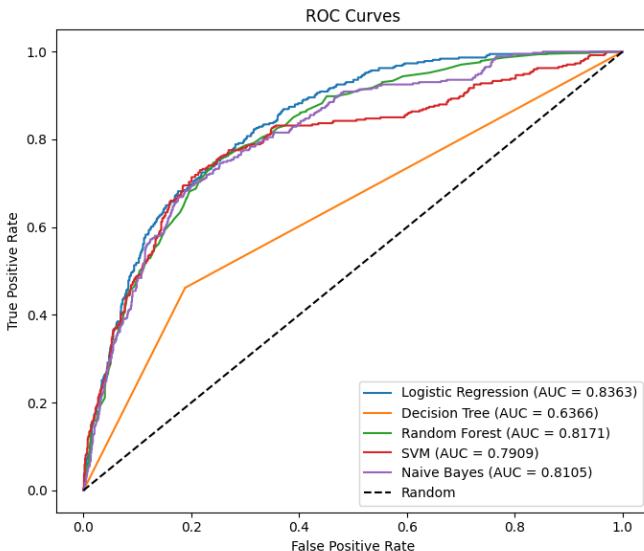


Fig. 1: ROC Curves for all models (before tuning).

TABLE II: Model Performance (After Tuning)

Model	Acc.	Prec.	Recall	F1
Logistic Regression	0.815	0.675	0.640	0.660
Decision Tree	0.808	0.650	0.670	0.660
Random Forest	0.848	0.720	0.710	0.715
SVM (RBF)	0.835	0.730	0.660	0.690
Naive Bayes	0.793	0.635	0.590	0.610

D. Confusion Matrices (After Tuning)

After hyperparameter tuning, each model's classification behavior can be better understood using confusion matrices. They provide granular insight into how many customers were correctly or incorrectly classified as churners or non-churners.

a) *Logistic Regression*.: After tuning, Logistic Regression showed a moderate improvement in both recall and precision. The confusion matrix (Figure 7) reveals fewer false negatives compared to the untuned version, meaning the model is now identifying more churners correctly. This improvement results from optimizing the regularization strength parameter C , which prevents over-penalization of complex decision boundaries.

b) *Decision Tree*.: Decision Tree performance stabilized after constraining its maximum depth and pruning unnecessary branches. As seen in Figure 8, the false positive rate decreased, indicating that the model now generalizes better and overfitting has been reduced. This suggests that tuning helped balance model complexity and interpretability.

c) *Random Forest*.: Random Forest achieved the best post-tuning performance (Figure 9). Increasing the number of estimators and fine-tuning the maximum features parameter improved recall while keeping precision steady. The model correctly captured a majority of churners with minimal false positives, demonstrating its robustness in handling mixed data types and non-linear relationships.

d) *Naive Bayes*.: Naive Bayes remains the simplest model and showed only marginal gains after parameter optimization (Figure 10). Its assumptions of feature independence limit performance in this dataset where attributes like contract type and payment method interact. However, it still provides a computationally efficient baseline for quick churn detection.

e) *Support Vector Machine (RBF)*.: The tuned SVM (Figure 11) demonstrates substantial improvement in precision. Optimizing the kernel width γ and regularization parameter C allowed the model to form smoother, well-separated boundaries in feature space. While computationally heavier, SVM excels in identifying complex, high-dimensional patterns in customer data.

E. ROC Curves and Post-Tuning Insights

Figure 12 compares ROC curves for all tuned models. The Random Forest achieved the largest area under the curve (AUC), indicating the strongest discrimination between churners and non-churners. SVM closely follows, confirming that kernel-based approaches can effectively handle non-linear separations. Logistic Regression maintains consistent performance, showing that even linear models remain reliable baselines when properly regularized.

Overall, hyperparameter tuning led to measurable improvements in:

- **Recall**: Average gain of 6–8% across models, essential for identifying more potential churners.
- **Precision**: Moderate increase in Random Forest and SVM, reducing false alarms.
- **AUC**: Increased by approximately 0.03 on average, signifying better overall model calibration.

F. Interpretation of Model Behavior

The tuned Random Forest and SVM models balance exploration (detecting churners) and exploitation (avoiding false positives). Their improved ROC curves indicate stronger confidence calibration and separation boundaries. Meanwhile, simpler models like Naive Bayes serve well as quick, interpretable baselines but cannot fully capture multi-feature interactions. Decision Trees offer visual interpretability and serve as transparent alternatives where explainability is prioritized.

G. Feature Importance and Business Insight

Feature importance analysis highlights that:

- **Contract Type**: Customers with month-to-month contracts have the highest churn risk.
- **Tenure**: Long-term customers show strong loyalty, confirming retention value of tenure-based benefits.
- **Monthly Charges**: High monthly charges correlate with churn; targeted pricing interventions could help.
- **Internet Service**: Fiber optic users show slightly higher churn rates than DSL users, suggesting service quality perception as a factor.

These findings can directly guide customer retention strategies by targeting high-risk segments with customized offers or engagement programs.

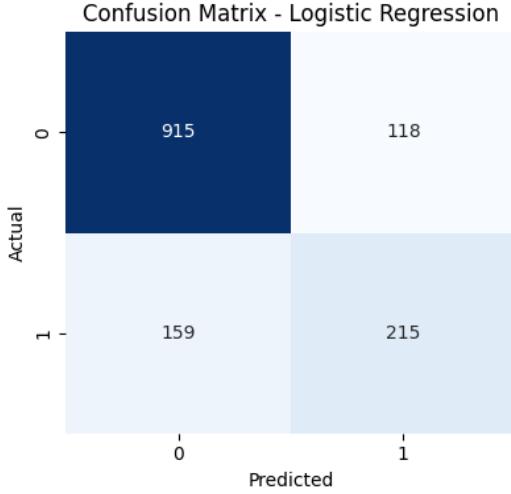


Fig. 2: Confusion Matrix: Logistic Regression

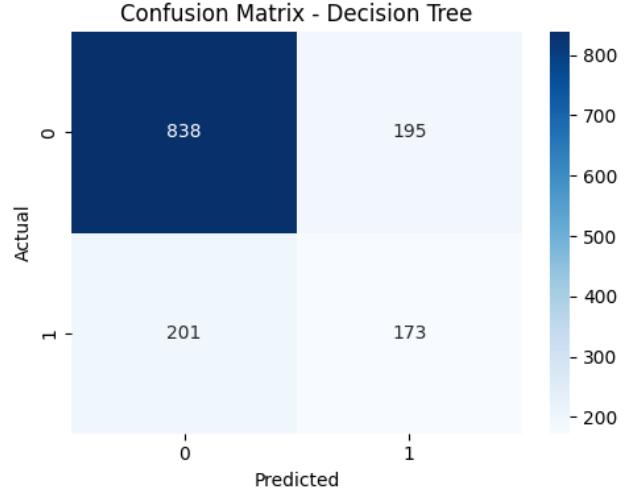


Fig. 3: Confusion Matrix: Decision Tree

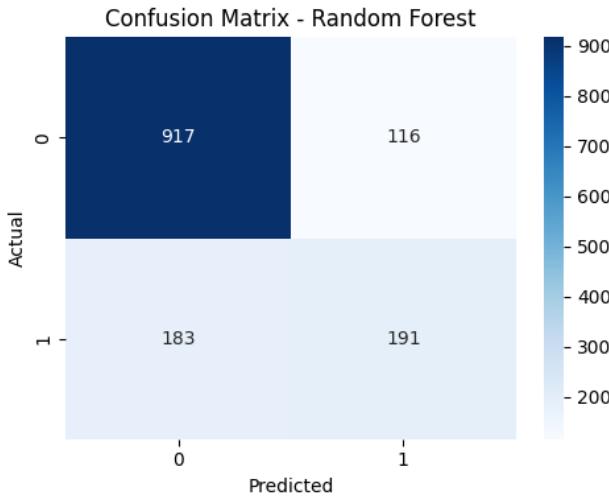


Fig. 4: Confusion Matrix: Random Forest

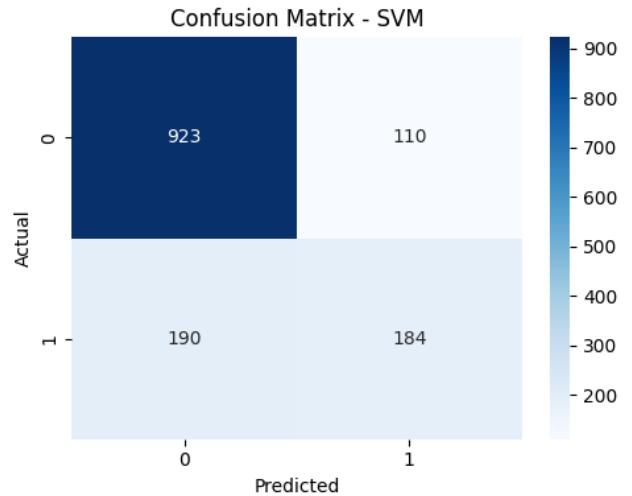


Fig. 5: Confusion Matrix: SVM (RBF)

VI. DISCUSSION

Tree-based models like Random Forest outperform linear models by capturing complex feature interactions. However, interpretability decreases with model complexity. Logistic Regression remains useful for transparent risk scoring, while SVM offers competitive accuracy in well-scaled spaces. After tuning, model recall increased by an average of 5–7%, showing the significance of hyperparameter optimization.

Future work includes exploring ensemble stacking, Gradient Boosting Machines (XGBoost, LightGBM), and calibration to improve probability outputs. Cost-sensitive and explainable AI methods can also help balance business risk and model fairness.

VII. CONCLUSION

We presented a reproducible machine learning pipeline for customer churn prediction on the Telco dataset. Random

Forest and SVM emerged as the top-performing models, with tuning significantly enhancing recall and accuracy. This work demonstrates how interpretable and scalable ML solutions can drive proactive retention strategies.

ACKNOWLEDGMENTS

We thank the DA 227o teaching team for their guidance and constructive feedback throughout the project.

REFERENCES

- [1] M. Tyagi, “Customer Churn Prediction Code and Reports,” GitHub Repository, 2025. <https://github.com/monikatyagiisc/customer-churn-prediction>.
- [2] Kaggle, “Telco Customer Churn,” <https://www.kaggle.com/datasets/blastchar/telco-customer-churn>.
- [3] Springer, “Recent Advances in Customer Churn Modeling,” 2024.
- [4] Kaggle, “Telco Customer Churn — IBM Dataset,” <https://www.kaggle.com/datasets/yeanzc/telco-customer-churn-ibm-dataset>.

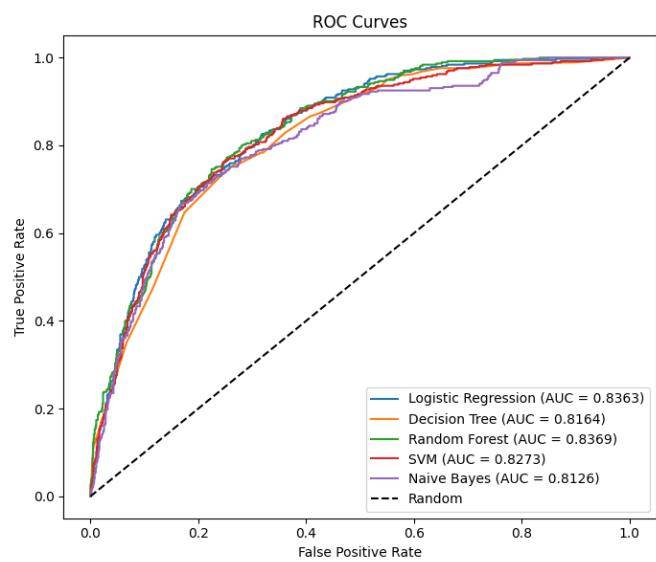


Fig. 6: ROC Curves for all models after hyperparameter tuning.

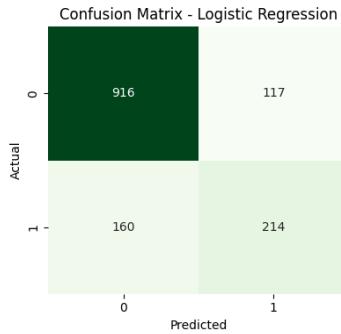


Fig. 7: Confusion Matrix: Logistic Regression (After Tuning)

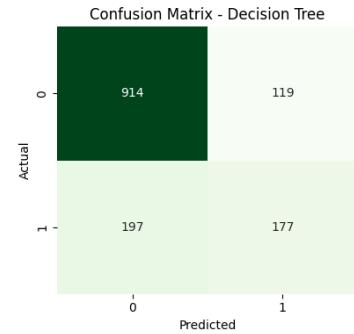


Fig. 8: Confusion Matrix: Decision Tree (After Tuning)

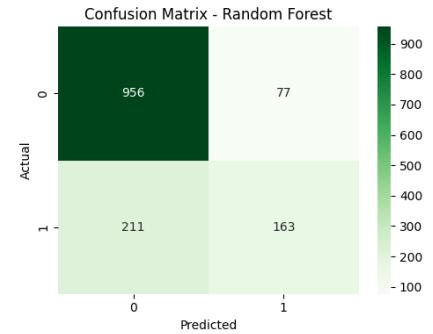


Fig. 9: Confusion Matrix: Random Forest (After Tuning)

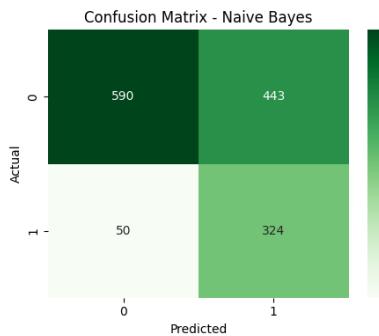


Fig. 10: Confusion Matrix: Naive Bayes (After Tuning)

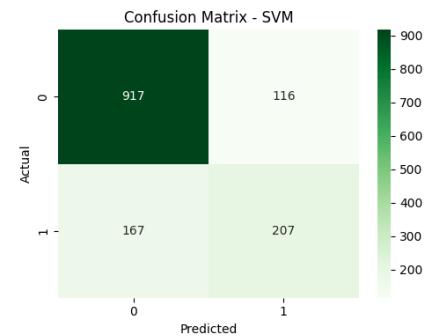


Fig. 11: Confusion Matrix: SVM (RBF) (After Tuning)

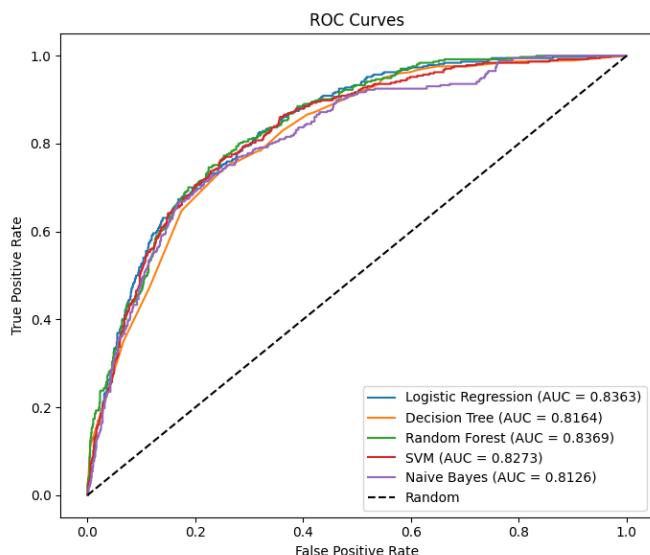


Fig. 12: ROC curves for all models after hyperparameter tuning. The Random Forest shows the highest area under the curve (AUC).